

# ***Animal Classification using Convolutional Neural Networks :-***

## *Abstract:*

This towering project is the enhancement of a Convolutional Neural Network (CNN) for animal images in fifteen types of animals. The only one-of-a-kind dataset having this smaller number of images is suited for beginner deep-learning applications. The images were pre-processed using TensorFlow utilities along with image augmentations that helped to minimize overfitting and better generalization of the model. The CNN architecture was made in Keras Sequential API. This CNN architecture was much simpler, but with notable effectiveness, composed of convolution, pooling, dropout, and dense layers. With approximately 96% training accuracy achieved and 98%+ validation accuracy barely requiring any GPU acceleration.

## *Introduction:*

- Image classification is a very basic, but very important, task in computer vision discriminative categorizing of objects presented in an image. When deep learning became in the limelight, CNNs were considered just the right tools: presenting them with raw-pixel data, they would then extract and learn the appropriate features.
- With this task, having to build an image classifier based on CNNs to identify animals in a dataset containing fifteen classes-dogs, cats, lions, tigers, elephants, and so on-is enough. Since the dataset is smaller and well-structured, it forms a great motivation for any beginner trying to understand and practice deep learning concepts.
- The model embraced TensorFlow and Keras, following an almost naive, working architecture for CNN modelling-conv, pool, dropout, and dense layers. Data augmentation methods like random flipping, rotation, and zooming were considered to improve generalization.
- This whole pipeline is lightweight enough to run on normal CPU hardware without any GPU acceleration requirements. The main motivation here is to demonstrate that the simplest kind of neural network can yield impressive accuracy on well-prepared datasets.

## *Dataset Description:*

- The dataset included 1,944 images divided into 15 animal categories: bear, bird, cat, cow, deer, dog, dolphin, elephant, giraffe, horse, kangaroo, lion, panda, tiger, and zebra. The TensorFlow tools automatically assign labels based on the class folders when loading data.
- All images were resized to 128×128 pixels. This keeps the size and shape consistent, which speeds up processing time. Since the dataset is small and balanced, it works well on computers that don't support GPUs.
- The `image_dataset_from_directory` function allows you to split the data into 80% for training and 20% for validation. It also handles batching and one-hot label encoding. Random flipping, rotation, zoom, and contrast adjustments improve the model's generalization.
- This well-organized and compact dataset provides a solid foundation for implementing and testing convolutional neural networks for multi-class image classification.

### *Technologies and Libraries Used:*

- So, here's the rundown: I built this whole thing in Python—because, let's be real, what else would you use for machine learning these days? Everything happened inside a Jupyter Notebook. If you've ever used one, you know it's basically the Swiss Army knife for coding, plotting, and tinkering without losing your mind.
- For the heavy lifting, I stuck with TensorFlow 2.x and its buddy, the good old Keras API. Keras makes building and training deep learning models feel almost... fun? Or at least less soul-crushing. Here's what I leaned on:
  1. TensorFlow / Keras: Built the CNN, handled data, did all the training jazz.
  2. NumPy: Crunches numbers, wrangles arrays—classic stuff.
  3. Matplotlib: Plots, charts, pretty graphs. Gotta see if your model's actually learning or just napping.
  4. OS: Moved files around, kept my folders from turning into digital spaghetti.
- Honestly, I ran the whole thing on a regular CPU laptop. No fancy GPUs, no server farm humming in the background—just a humble machine and a small dataset. Means anyone can try it, even if you're not swimming in hardware.
- All these pieces came together to make a clean, modular setup for building and testing an animal classifier. Nothing convoluted. Nothing locked behind a \$10,000 graphics card. Just straight-up code and some patience.

## Data Preprocessing:

- Alright, so here's how I wrangled the data before tossing it into the model. First off, TensorFlow's got this handy `image_dataset_from_directory()` thing—super clutch. You just point it at your folders, and boom, it figures out your labels, squishes everything to 128x128, and even does the one-hot label jazz for you. No more manual label headaches.
- Next, I chopped up the data: 80% for training, 20% for validation. Gotta keep that model honest, right? Also, shuffled the batches—don't want the thing picking up on some weird order and thinking that's important.
- Now, let's talk about data augmentation. Because, honestly, if you don't mess with your images a bit, you're just asking for overfitting. I tossed in some random flips, spins, zooms, and even fiddled with the contrast. All that chaos only happens during training, though—validation data stays squeaky clean. No cheating.
- Oh, and I made sure all the pixel values landed between 0 and 1 (thank you, `Rescaling(1/255)`). That makes the training smoother and stops things from exploding numerically.
- So yeah, without all these steps, you're basically asking your model to learn from a hot mess. This way, it actually finds the good stuff and generalizes, instead of just memorizing the training set like a parrot.

## Model Architecture:

The image bracket model was erected using a successional Convolutional Neural Network (CNN) armature enforced via Keras. The network is designed to prize spatial scales from images and learn meaningful features for accurate beast bracket. The armature consists of the following layers:

- Data Augmentation Layers Applied arbitrary flips, reels, thrums, and discrepancy adaptations to increase training data variability.
- Rescaling Subcaste regularized image pixel values from (0, 255) to (0, 1) for briskly and more stable training.
- Convolutional Layers Three Conv2D layers with adding sludge sizes (e.g., 32, 64, 128), each followed by a ReLU activation and MaxPooling2D to reduce spatial confines.
- Powerhouse Layers Added between complication blocks to help overfitting by aimlessly disabling neurons during training.
- Flatten Layer Converted the 2D point maps into a 1D vector.

- Completely Connected (thick) Layers Included one or further thick layers with ReLU activation to learn complex patterns.
- Affair Subcaste A thick subcaste with 15 units (for each class) and a Softmax activation to produce class chances.

This armature was featherlight yet important enough to achieve high delicacy on the small dataset without taking GPU coffers.

### Training Configuration:

The model was trained using the TensorFlow Keras API, with settings chosen to balance training effectiveness and delicacy. The following configuration was used:

- Loss Function CategoricalCrossentropy, suitable for multi-class bracket with one-hot decoded markers.
- Optimizer Adam, an adaptive literacy rate optimizer known for fast confluence and stability.
- Evaluation Metric Accuracy, to cover the proportion of rightly classified images.
- Batch Size 32 images per batch.
- Ages 200 training cycles, which handed sufficient time for confluence without overfitting.

The training process was run entirely on a CPU, as the dataset was small and the model featherlight. Data addition and rescaling were applied in real- time during each time using Keras preprocessing layers.

Confirmation delicacy and loss were covered after each time to estimate conception performance. No beforehand stopping or learning rate scheduling was needed due to the model's harmonious confluence.


This configuration proved to be effective, achieving high training and confirmation delicacy with smooth loss angles, indicating stable and effective training behaviour throughout the process.

### Results and Evaluation:

- The model performed very well on both the training and the evidence sets, indicating that it could have generalized from a rather small dataset.
- Final Training accuracy: 96
- Final evidence accuracy: 98

- In the course of the 200 training sessions, the loss for evidence-training and training datasets consistently decreased and the sensitivity gradually grew. There was no evidence of material overfitting in the models, due to the inclusion of hustler layers and data augmentation paths.
- Lineplots learned accuracy and loss of other criteria across epochs. These plots confirmed smooth convergence, as the evidence angles only slightly deviated from the training angles.
- Sample prognostications were also evaluated on evidence images, and the model accurately predicted most images, including those with different backgrounds and lightening.
- Although we did not make use of a confusion matrix, visual inspection of the predictions revealed that the position of class- position delicacy was high for all higher orders.
- In conclusion, this study confirms the efficiency of a well- tuned CNN architecture for multi- class image type based on feathery resources. The model is robust and can be applied to an extend-to-real-time or push- to-deploy operation.

### Results Graph:



### Sample Output:

Prediction: Bear (99.91%)



---

### Challenges Faced:

Although the dataset and model were fairly simple, many challenges arose during the development process:

#### 1. Class Imbalance:

Some beast classes had slightly smaller images than others, which needed careful observation to ensure that the model was not prejudiced toward further frequent classes.

#### 2. Overfitting threat:

Due to the small dataset size, the model snappily began to study training images. This was eased by introducing powerhouse layers and applying data addition to instinctively increase data diversity.

#### 3. Augmentation Tuning:

Chancing the right balance in addition parameters was pivotal. Too important deformation affected the image quality and hurt delicacy, while too little had no effect on conception.

#### 4. Model Complexity:

Originally, deeper models were tested but redounded in slower training without significant delicacy earnings. The armature was optimized to be featherlight and effective for CPU training.

#### 5. Confirmation Monitoring:

Without using calls like

EarlyStopping, covering training manually was necessary to ensure the model did not begin to overfit in after ages.

These challenges handed precious literacy gestures and guided architectural and preprocessing opinions that led to a stable, high- performing model.

### Future Scope:

- Though the present model does achieve great results on a small and clean dataset, there could be many improvements and extensions to the present project:
- Larger & More Diverse Data: Extending our data by including more animals or higher diversity in images (backgrounds, illumination, views) would evaluate how well the model would generalize and how robust it is.
- Transfer Learning: Using pretrained models will definitely help out (e.g., VGG16, ResNet, or MobileNet) especially when dealing with large datasets -- :) it will help to increase the accuracy and save time for training.
- Real-Time Deployment: Embed the trained model to a real-time application (running on either mobile app or web application) using TensorFlow Lite or TensorFlow.js size and inference time can be reduced using the techniques such as quantization or pruning, which can be helpful to make it work efficiently on edge devices or low-resource environments.
- Explainable AI (XAI): By employing models such as Grad-CAM or LIME, we can understand which areas of an image are being considered by CNNs, making model decisions transparent and trustworthy.



- Confusion Matrix and Metrics: Knowledge of a confusion matrix, precision, recall, and F1-score, would give more insight on how the model is doing per class, showing where it may apply more effort.

## Conclusions:

This project effectively proves the strength and applicability of CNN in animal multi-class image classification with a small and clear animal dataset. Despite small dataset and being trained on a normal machine with limited power, the model had an impressive accuracy, confirming that with the correct preprocessing, data augmentation, and well-tuned architecture, a good performance can be achieved without GPU support.

The experiment also demonstrated the similar challenges like class imbalance, overfitting which was successfully handled with dropout layers and weak augmentation. Finally, this work demonstrates that even a small CNN can be meaningful, and it provides a strong basis for larger, scalable, deployable and interpretable models.



```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Dropout, Rescaling
from tensorflow.keras.layers import RandomFlip, RandomRotation,
RandomZoom, RandomContrast
import numpy as np
import matplotlib.pyplot as plt
import os
```

```
DATA_DIR = 'E:\\Projects\\animal_classification\\Animal
Classification\\dataset'
```

```
IMAGE_SIZE = (128, 128)
BATCH_SIZE = 32
VALIDATION_SPLIT = 0.2
```

```
train_ds = tf.keras.utils.image_dataset_from_directory(
    DATA_DIR,
    labels='inferred',
    label_mode='categorical',
    image_size=IMAGE_SIZE,
    interpolation='nearest',
    batch_size=BATCH_SIZE,
    shuffle=True,
    validation_split=VALIDATION_SPLIT,
    subset='training',
    seed=123
)
```

```
val_ds = tf.keras.utils.image_dataset_from_directory(
    DATA_DIR,
    labels='inferred',
    label_mode='categorical',
    image_size=IMAGE_SIZE,
    interpolation='nearest',
    batch_size=BATCH_SIZE,
    shuffle=False,
    validation_split=VALIDATION_SPLIT,
    subset='validation',
    seed=123
)
```

```
Found 1944 files belonging to 15 classes.
Using 1556 files for training.
Found 1944 files belonging to 15 classes.
Using 388 files for validation.
```

```
class_names = train_ds.class_names
NUM_CLASSES = len(class_names)
```

```

print(f"Detected {NUM_CLASSES} classes: {class_names}")

AUTOTUNE = tf.data.AUTOTUNE
train_ds =
train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

data_augmentation = Sequential([
    RandomFlip("horizontal_and_vertical"),
    RandomRotation(0.2),
    RandomZoom(0.2),
    RandomContrast(0.2),
], name="data_augmentation")

Detected 15 classes: ['Bear', 'Bird', 'Cat', 'Cow', 'Deer', 'Dog',
'Dolphin', 'Elephant', 'Giraffe', 'Horse', 'Kangaroo', 'Lion',
'Panda', 'Tiger', 'Zebra']

model = Sequential([
    data_augmentation,
    Rescaling(1./255),
    Conv2D(32, (3, 3), activation='relu', input_shape=(IMAGE_SIZE[0],
IMAGE_SIZE[1], 3)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(NUM_CLASSES, activation='softmax')
])

model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

model.summary()

EPOCHS = 200

history = model.fit(
    train_ds,
    epochs=EPOCHS,
    validation_data=val_ds,
)

```

```

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(EPOCHS)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

model.save('animal_classifier.keras')
print("Model saved as 'animal_classifier.keras'")

```

```

C:\Users\tanis\AppData\Local\Packages\
PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-
packages\Python311\site-packages\keras\src\layers\convolutional\
base_conv.py:99: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(

```

Model: "sequential"

Layer (type) Param #	Output Shape	
data_augmentation (Sequential) (unbuilt)	?	0
rescaling (Rescaling) (unbuilt)	?	0
conv2d (Conv2D) (unbuilt)	?	0

max_pooling2d (MaxPooling2D)	?	0
(unbuilt)		
conv2d_1 (Conv2D)	?	0
(unbuilt)		
max_pooling2d_1 (MaxPooling2D)	?	0
(unbuilt)		
conv2d_2 (Conv2D)	?	0
(unbuilt)		
max_pooling2d_2 (MaxPooling2D)	?	0
(unbuilt)		
flatten (Flatten)	?	0
(unbuilt)		
dense (Dense)	?	0
(unbuilt)		
dropout (Dropout)	?	
0		
dense_1 (Dense)	?	0
(unbuilt)		

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Epoch 1/200

49/49 ————— 50s 677ms/step - accuracy: 0.0960 - loss: 2.7392 - val\_accuracy: 0.1985 - val\_loss: 2.6428

Epoch 2/200

49/49 ————— 30s 613ms/step - accuracy: 0.1459 - loss:

2.5641 - val\_accuracy: 0.3660 - val\_loss: 2.0371  
Epoch 3/200  
49/49 \_\_\_\_\_ 29s 593ms/step - accuracy: 0.2087 - loss: 2.4206 - val\_accuracy: 0.6340 - val\_loss: 1.2463  
Epoch 4/200  
49/49 \_\_\_\_\_ 30s 620ms/step - accuracy: 0.2715 - loss: 2.2607 - val\_accuracy: 0.4974 - val\_loss: 1.5235  
Epoch 5/200  
49/49 \_\_\_\_\_ 29s 596ms/step - accuracy: 0.2912 - loss: 2.1497 - val\_accuracy: 0.7526 - val\_loss: 0.9594  
Epoch 6/200  
49/49 \_\_\_\_\_ 29s 594ms/step - accuracy: 0.3487 - loss: 2.0438 - val\_accuracy: 0.6727 - val\_loss: 1.3251  
Epoch 7/200  
49/49 \_\_\_\_\_ 30s 604ms/step - accuracy: 0.3546 - loss: 1.9902 - val\_accuracy: 0.7062 - val\_loss: 1.3023  
Epoch 8/200  
49/49 \_\_\_\_\_ 29s 600ms/step - accuracy: 0.3525 - loss: 1.9313 - val\_accuracy: 0.7139 - val\_loss: 0.9221  
Epoch 9/200  
49/49 \_\_\_\_\_ 30s 616ms/step - accuracy: 0.4037 - loss: 1.8503 - val\_accuracy: 0.8454 - val\_loss: 0.6770  
Epoch 10/200  
49/49 \_\_\_\_\_ 30s 607ms/step - accuracy: 0.4364 - loss: 1.7581 - val\_accuracy: 0.6392 - val\_loss: 1.4580  
Epoch 11/200  
49/49 \_\_\_\_\_ 30s 617ms/step - accuracy: 0.4165 - loss: 1.7744 - val\_accuracy: 0.6392 - val\_loss: 1.2209  
Epoch 12/200  
49/49 \_\_\_\_\_ 29s 598ms/step - accuracy: 0.4505 - loss: 1.7574 - val\_accuracy: 0.8299 - val\_loss: 0.5701  
Epoch 13/200  
49/49 \_\_\_\_\_ 29s 598ms/step - accuracy: 0.4887 - loss: 1.5600 - val\_accuracy: 0.7036 - val\_loss: 0.9373  
Epoch 14/200  
49/49 \_\_\_\_\_ 29s 599ms/step - accuracy: 0.4793 - loss: 1.5720 - val\_accuracy: 0.8093 - val\_loss: 0.6839  
Epoch 15/200  
49/49 \_\_\_\_\_ 30s 623ms/step - accuracy: 0.5293 - loss: 1.4714 - val\_accuracy: 0.7887 - val\_loss: 0.6839  
Epoch 16/200  
49/49 \_\_\_\_\_ 29s 598ms/step - accuracy: 0.5402 - loss: 1.4363 - val\_accuracy: 0.7732 - val\_loss: 0.7526  
Epoch 17/200  
49/49 \_\_\_\_\_ 29s 593ms/step - accuracy: 0.5105 - loss: 1.5123 - val\_accuracy: 0.7552 - val\_loss: 0.7866  
Epoch 18/200  
49/49 \_\_\_\_\_ 29s 603ms/step - accuracy: 0.5557 - loss: 1.4071 - val\_accuracy: 0.6804 - val\_loss: 1.8622

Epoch 19/200  
49/49 \_\_\_\_\_ 30s 616ms/step - accuracy: 0.5715 - loss: 1.3324 - val\_accuracy: 0.7242 - val\_loss: 1.3533  
Epoch 20/200  
49/49 \_\_\_\_\_ 30s 621ms/step - accuracy: 0.5913 - loss: 1.2396 - val\_accuracy: 0.7655 - val\_loss: 0.9209  
Epoch 21/200  
49/49 \_\_\_\_\_ 29s 601ms/step - accuracy: 0.5773 - loss: 1.2636 - val\_accuracy: 0.8144 - val\_loss: 0.8158  
Epoch 22/200  
49/49 \_\_\_\_\_ 30s 603ms/step - accuracy: 0.5867 - loss: 1.2482 - val\_accuracy: 0.7345 - val\_loss: 1.4049  
Epoch 23/200  
49/49 \_\_\_\_\_ 29s 602ms/step - accuracy: 0.6172 - loss: 1.1522 - val\_accuracy: 0.8969 - val\_loss: 0.4576  
Epoch 24/200  
49/49 \_\_\_\_\_ 30s 609ms/step - accuracy: 0.5779 - loss: 1.3087 - val\_accuracy: 0.8686 - val\_loss: 0.5080  
Epoch 25/200  
49/49 \_\_\_\_\_ 30s 621ms/step - accuracy: 0.6155 - loss: 1.0873 - val\_accuracy: 0.8351 - val\_loss: 0.6127  
Epoch 26/200  
49/49 \_\_\_\_\_ 30s 618ms/step - accuracy: 0.6257 - loss: 1.1019 - val\_accuracy: 0.8814 - val\_loss: 0.5174  
Epoch 27/200  
49/49 \_\_\_\_\_ 29s 595ms/step - accuracy: 0.6621 - loss: 1.0790 - val\_accuracy: 0.7809 - val\_loss: 0.9482  
Epoch 28/200  
49/49 \_\_\_\_\_ 30s 604ms/step - accuracy: 0.6217 - loss: 1.0910 - val\_accuracy: 0.7397 - val\_loss: 1.9697  
Epoch 29/200  
49/49 \_\_\_\_\_ 29s 599ms/step - accuracy: 0.6474 - loss: 1.0548 - val\_accuracy: 0.7629 - val\_loss: 1.2100  
Epoch 30/200  
49/49 \_\_\_\_\_ 29s 595ms/step - accuracy: 0.6496 - loss: 1.0208 - val\_accuracy: 0.7655 - val\_loss: 1.3318  
Epoch 31/200  
49/49 \_\_\_\_\_ 30s 608ms/step - accuracy: 0.6585 - loss: 1.0008 - val\_accuracy: 0.8299 - val\_loss: 0.7257  
Epoch 32/200  
49/49 \_\_\_\_\_ 29s 594ms/step - accuracy: 0.6796 - loss: 0.9041 - val\_accuracy: 0.8428 - val\_loss: 0.6770  
Epoch 33/200  
49/49 \_\_\_\_\_ 29s 593ms/step - accuracy: 0.6869 - loss: 0.9538 - val\_accuracy: 0.8557 - val\_loss: 0.7184  
Epoch 34/200  
49/49 \_\_\_\_\_ 29s 599ms/step - accuracy: 0.6912 - loss: 0.8939 - val\_accuracy: 0.8196 - val\_loss: 0.6942  
Epoch 35/200

49/49 ————— 29s 595ms/step - accuracy: 0.6736 - loss: 0.9369 - val\_accuracy: 0.8376 - val\_loss: 0.8627  
Epoch 36/200  
49/49 ————— 30s 612ms/step - accuracy: 0.7009 - loss: 0.9027 - val\_accuracy: 0.9046 - val\_loss: 0.4854  
Epoch 37/200  
49/49 ————— 30s 618ms/step - accuracy: 0.7031 - loss: 0.8443 - val\_accuracy: 0.8170 - val\_loss: 0.9608  
Epoch 38/200  
49/49 ————— 29s 601ms/step - accuracy: 0.7177 - loss: 0.8645 - val\_accuracy: 0.8660 - val\_loss: 0.6042  
Epoch 39/200  
49/49 ————— 29s 592ms/step - accuracy: 0.7642 - loss: 0.7395 - val\_accuracy: 0.9536 - val\_loss: 0.1506  
Epoch 40/200  
49/49 ————— 29s 600ms/step - accuracy: 0.7470 - loss: 0.7856 - val\_accuracy: 0.7603 - val\_loss: 1.4716  
Epoch 41/200  
49/49 ————— 30s 605ms/step - accuracy: 0.7294 - loss: 0.7852 - val\_accuracy: 0.7990 - val\_loss: 1.2621  
Epoch 42/200  
49/49 ————— 30s 616ms/step - accuracy: 0.7246 - loss: 0.8221 - val\_accuracy: 0.8608 - val\_loss: 0.5445  
Epoch 43/200  
49/49 ————— 29s 599ms/step - accuracy: 0.7427 - loss: 0.7126 - val\_accuracy: 0.8119 - val\_loss: 0.9179  
Epoch 44/200  
49/49 ————— 29s 600ms/step - accuracy: 0.7665 - loss: 0.6755 - val\_accuracy: 0.8660 - val\_loss: 0.7541  
Epoch 45/200  
49/49 ————— 30s 603ms/step - accuracy: 0.7472 - loss: 0.7470 - val\_accuracy: 0.8814 - val\_loss: 0.5126  
Epoch 46/200  
49/49 ————— 30s 604ms/step - accuracy: 0.7235 - loss: 0.8407 - val\_accuracy: 0.8711 - val\_loss: 0.6504  
Epoch 47/200  
49/49 ————— 30s 611ms/step - accuracy: 0.7548 - loss: 0.7020 - val\_accuracy: 0.9175 - val\_loss: 0.3453  
Epoch 48/200  
49/49 ————— 31s 623ms/step - accuracy: 0.7787 - loss: 0.6798 - val\_accuracy: 0.9098 - val\_loss: 0.3581  
Epoch 49/200  
49/49 ————— 29s 598ms/step - accuracy: 0.8010 - loss: 0.6008 - val\_accuracy: 0.8428 - val\_loss: 0.9900  
Epoch 50/200  
49/49 ————— 29s 598ms/step - accuracy: 0.7849 - loss: 0.6697 - val\_accuracy: 0.9407 - val\_loss: 0.3188  
Epoch 51/200  
49/49 ————— 29s 597ms/step - accuracy: 0.7717 - loss:



0.6830 - val\_accuracy: 0.9098 - val\_loss: 0.3688  
Epoch 52/200  
49/49 \_\_\_\_\_ 29s 597ms/step - accuracy: 0.7997 - loss:  
0.6036 - val\_accuracy: 0.8763 - val\_loss: 0.5629  
Epoch 53/200  
49/49 \_\_\_\_\_ 29s 599ms/step - accuracy: 0.8020 - loss:  
0.6019 - val\_accuracy: 0.8969 - val\_loss: 0.6184  
Epoch 54/200  
49/49 \_\_\_\_\_ 30s 616ms/step - accuracy: 0.7550 - loss:  
0.7506 - val\_accuracy: 0.8943 - val\_loss: 0.7691  
Epoch 55/200  
49/49 \_\_\_\_\_ 24s 495ms/step - accuracy: 0.7887 - loss:  
0.6044 - val\_accuracy: 0.8814 - val\_loss: 0.5895  
Epoch 56/200  
49/49 \_\_\_\_\_ 16s 337ms/step - accuracy: 0.8080 - loss:  
0.5545 - val\_accuracy: 0.8918 - val\_loss: 0.6723  
Epoch 57/200  
49/49 \_\_\_\_\_ 17s 341ms/step - accuracy: 0.8092 - loss:  
0.6214 - val\_accuracy: 0.9072 - val\_loss: 0.3825  
Epoch 58/200  
49/49 \_\_\_\_\_ 17s 338ms/step - accuracy: 0.8110 - loss:  
0.5531 - val\_accuracy: 0.8789 - val\_loss: 0.5847  
Epoch 59/200  
49/49 \_\_\_\_\_ 17s 337ms/step - accuracy: 0.7992 - loss:  
0.6265 - val\_accuracy: 0.9227 - val\_loss: 0.4595  
Epoch 60/200  
49/49 \_\_\_\_\_ 17s 349ms/step - accuracy: 0.8050 - loss:  
0.6373 - val\_accuracy: 0.8067 - val\_loss: 1.9501  
Epoch 61/200  
49/49 \_\_\_\_\_ 17s 337ms/step - accuracy: 0.8235 - loss:  
0.5564 - val\_accuracy: 0.8376 - val\_loss: 1.3702  
Epoch 62/200  
49/49 \_\_\_\_\_ 16s 332ms/step - accuracy: 0.8284 - loss:  
0.4982 - val\_accuracy: 0.9124 - val\_loss: 0.5532  
Epoch 63/200  
49/49 \_\_\_\_\_ 16s 335ms/step - accuracy: 0.8192 - loss:  
0.5469 - val\_accuracy: 0.8686 - val\_loss: 1.0538  
Epoch 64/200  
49/49 \_\_\_\_\_ 16s 336ms/step - accuracy: 0.8459 - loss:  
0.4624 - val\_accuracy: 0.9227 - val\_loss: 0.5016  
Epoch 65/200  
49/49 \_\_\_\_\_ 17s 346ms/step - accuracy: 0.8384 - loss:  
0.5062 - val\_accuracy: 0.9381 - val\_loss: 0.2818  
Epoch 66/200  
49/49 \_\_\_\_\_ 18s 369ms/step - accuracy: 0.8394 - loss:  
0.4945 - val\_accuracy: 0.8582 - val\_loss: 0.9490  
Epoch 67/200  
49/49 \_\_\_\_\_ 18s 363ms/step - accuracy: 0.8394 - loss:  
0.4493 - val\_accuracy: 0.9201 - val\_loss: 0.3892

Epoch 68/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.7809 - loss: 0.6760 - val\_accuracy: 0.8943 - val\_loss: 0.6049  
Epoch 69/200  
49/49 \_\_\_\_\_ 17s 346ms/step - accuracy: 0.8060 - loss: 0.6029 - val\_accuracy: 0.8351 - val\_loss: 0.8123  
Epoch 70/200  
49/49 \_\_\_\_\_ 17s 338ms/step - accuracy: 0.8656 - loss: 0.4188 - val\_accuracy: 0.9485 - val\_loss: 0.2256  
Epoch 71/200  
49/49 \_\_\_\_\_ 17s 342ms/step - accuracy: 0.8448 - loss: 0.4534 - val\_accuracy: 0.9356 - val\_loss: 0.3890  
Epoch 72/200  
49/49 \_\_\_\_\_ 17s 346ms/step - accuracy: 0.8585 - loss: 0.4424 - val\_accuracy: 0.9124 - val\_loss: 0.6312  
Epoch 73/200  
49/49 \_\_\_\_\_ 17s 339ms/step - accuracy: 0.8009 - loss: 0.6148 - val\_accuracy: 0.9381 - val\_loss: 0.3789  
Epoch 74/200  
49/49 \_\_\_\_\_ 17s 346ms/step - accuracy: 0.8421 - loss: 0.4483 - val\_accuracy: 0.9227 - val\_loss: 0.6145  
Epoch 75/200  
49/49 \_\_\_\_\_ 17s 350ms/step - accuracy: 0.8368 - loss: 0.4913 - val\_accuracy: 0.9897 - val\_loss: 0.0404  
Epoch 76/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.8279 - loss: 0.5393 - val\_accuracy: 0.9021 - val\_loss: 0.5877  
Epoch 77/200  
49/49 \_\_\_\_\_ 17s 351ms/step - accuracy: 0.8489 - loss: 0.4437 - val\_accuracy: 0.9716 - val\_loss: 0.1200  
Epoch 78/200  
49/49 \_\_\_\_\_ 18s 367ms/step - accuracy: 0.8519 - loss: 0.4646 - val\_accuracy: 0.9407 - val\_loss: 0.4235  
Epoch 79/200  
49/49 \_\_\_\_\_ 18s 359ms/step - accuracy: 0.8616 - loss: 0.4052 - val\_accuracy: 0.9588 - val\_loss: 0.2722  
Epoch 80/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8395 - loss: 0.4490 - val\_accuracy: 0.9046 - val\_loss: 0.9092  
Epoch 81/200  
49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.8652 - loss: 0.4193 - val\_accuracy: 0.9613 - val\_loss: 0.2022  
Epoch 82/200  
49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.8302 - loss: 0.5001 - val\_accuracy: 0.8531 - val\_loss: 1.1137  
Epoch 83/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8475 - loss: 0.4845 - val\_accuracy: 0.9639 - val\_loss: 0.2739  
Epoch 84/200

49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.8780 - loss: 0.4210 - val\_accuracy: 0.8582 - val\_loss: 1.0792  
Epoch 85/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8817 - loss: 0.3920 - val\_accuracy: 0.9485 - val\_loss: 0.2965  
Epoch 86/200  
49/49 \_\_\_\_\_ 18s 369ms/step - accuracy: 0.8728 - loss: 0.4154 - val\_accuracy: 0.8222 - val\_loss: 1.6733  
Epoch 87/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8620 - loss: 0.3998 - val\_accuracy: 0.8660 - val\_loss: 0.7875  
Epoch 88/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8574 - loss: 0.4098 - val\_accuracy: 0.8454 - val\_loss: 1.0201  
Epoch 89/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8646 - loss: 0.3717 - val\_accuracy: 0.9588 - val\_loss: 0.1907  
Epoch 90/200  
49/49 \_\_\_\_\_ 17s 352ms/step - accuracy: 0.8816 - loss: 0.3980 - val\_accuracy: 0.8918 - val\_loss: 1.0449  
Epoch 91/200  
49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.8614 - loss: 0.4297 - val\_accuracy: 0.9253 - val\_loss: 0.6051  
Epoch 92/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.8545 - loss: 0.4457 - val\_accuracy: 0.8557 - val\_loss: 0.9849  
Epoch 93/200  
49/49 \_\_\_\_\_ 19s 392ms/step - accuracy: 0.8748 - loss: 0.3957 - val\_accuracy: 0.9716 - val\_loss: 0.0778  
Epoch 94/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8874 - loss: 0.3135 - val\_accuracy: 0.8918 - val\_loss: 0.8085  
Epoch 95/200  
49/49 \_\_\_\_\_ 18s 368ms/step - accuracy: 0.8697 - loss: 0.4037 - val\_accuracy: 0.9201 - val\_loss: 0.7278  
Epoch 96/200  
49/49 \_\_\_\_\_ 18s 357ms/step - accuracy: 0.8892 - loss: 0.3317 - val\_accuracy: 0.8737 - val\_loss: 1.3129  
Epoch 97/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8510 - loss: 0.4480 - val\_accuracy: 0.9948 - val\_loss: 0.0228  
Epoch 98/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8791 - loss: 0.3459 - val\_accuracy: 0.9768 - val\_loss: 0.0960  
Epoch 99/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9061 - loss: 0.3020 - val\_accuracy: 0.9407 - val\_loss: 0.4869  
Epoch 100/200  
49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.8838 - loss:

0.3276 - val\_accuracy: 0.9536 - val\_loss: 0.3760  
Epoch 101/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8896 - loss:  
0.3061 - val\_accuracy: 0.8840 - val\_loss: 1.1959  
Epoch 102/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8989 - loss:  
0.3093 - val\_accuracy: 0.8995 - val\_loss: 1.0291  
Epoch 103/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8876 - loss:  
0.3172 - val\_accuracy: 0.8582 - val\_loss: 1.4319  
Epoch 104/200  
49/49 \_\_\_\_\_ 17s 352ms/step - accuracy: 0.8934 - loss:  
0.3619 - val\_accuracy: 0.9356 - val\_loss: 0.7789  
Epoch 105/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.8825 - loss:  
0.3246 - val\_accuracy: 0.9665 - val\_loss: 0.1189  
Epoch 106/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8996 - loss:  
0.3251 - val\_accuracy: 0.9407 - val\_loss: 0.2137  
Epoch 107/200  
49/49 \_\_\_\_\_ 18s 362ms/step - accuracy: 0.8848 - loss:  
0.3677 - val\_accuracy: 0.9072 - val\_loss: 0.7277  
Epoch 108/200  
49/49 \_\_\_\_\_ 17s 352ms/step - accuracy: 0.8809 - loss:  
0.3995 - val\_accuracy: 0.9845 - val\_loss: 0.0387  
Epoch 109/200  
49/49 \_\_\_\_\_ 18s 359ms/step - accuracy: 0.8654 - loss:  
0.3919 - val\_accuracy: 0.9613 - val\_loss: 0.0863  
Epoch 110/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.8892 - loss:  
0.3642 - val\_accuracy: 0.9407 - val\_loss: 0.3160  
Epoch 111/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8992 - loss:  
0.3037 - val\_accuracy: 0.9897 - val\_loss: 0.0333  
Epoch 112/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8934 - loss:  
0.3187 - val\_accuracy: 0.9820 - val\_loss: 0.0622  
Epoch 113/200  
49/49 \_\_\_\_\_ 18s 368ms/step - accuracy: 0.9154 - loss:  
0.2671 - val\_accuracy: 0.9562 - val\_loss: 0.2333  
Epoch 114/200  
49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.9012 - loss:  
0.3509 - val\_accuracy: 0.9613 - val\_loss: 0.1826  
Epoch 115/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.8849 - loss:  
0.3758 - val\_accuracy: 0.9639 - val\_loss: 0.1904  
Epoch 116/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.8995 - loss:  
0.3152 - val\_accuracy: 0.9536 - val\_loss: 0.1457

Epoch 117/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8953 - loss: 0.2802 - val\_accuracy: 0.9227 - val\_loss: 0.5656  
Epoch 118/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.9065 - loss: 0.3016 - val\_accuracy: 0.9794 - val\_loss: 0.0353  
Epoch 119/200  
49/49 \_\_\_\_\_ 18s 359ms/step - accuracy: 0.8994 - loss: 0.3260 - val\_accuracy: 0.9948 - val\_loss: 0.0148  
Epoch 120/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.8815 - loss: 0.3624 - val\_accuracy: 0.9820 - val\_loss: 0.0666  
Epoch 121/200  
49/49 \_\_\_\_\_ 17s 351ms/step - accuracy: 0.8941 - loss: 0.3037 - val\_accuracy: 0.9768 - val\_loss: 0.1016  
Epoch 122/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8805 - loss: 0.3731 - val\_accuracy: 0.8737 - val\_loss: 1.0286  
Epoch 123/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8991 - loss: 0.3542 - val\_accuracy: 0.9562 - val\_loss: 0.1673  
Epoch 124/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.8975 - loss: 0.3114 - val\_accuracy: 0.9845 - val\_loss: 0.0435  
Epoch 125/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.8803 - loss: 0.3307 - val\_accuracy: 0.9588 - val\_loss: 0.1499  
Epoch 126/200  
49/49 \_\_\_\_\_ 18s 357ms/step - accuracy: 0.8882 - loss: 0.3494 - val\_accuracy: 0.9820 - val\_loss: 0.0572  
Epoch 127/200  
49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.9038 - loss: 0.2999 - val\_accuracy: 0.9588 - val\_loss: 0.3420  
Epoch 128/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9209 - loss: 0.2476 - val\_accuracy: 0.9588 - val\_loss: 0.1812  
Epoch 129/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.9280 - loss: 0.2308 - val\_accuracy: 0.9459 - val\_loss: 0.4237  
Epoch 130/200  
49/49 \_\_\_\_\_ 18s 363ms/step - accuracy: 0.9077 - loss: 0.2795 - val\_accuracy: 0.9227 - val\_loss: 0.5126  
Epoch 131/200  
49/49 \_\_\_\_\_ 18s 362ms/step - accuracy: 0.8978 - loss: 0.2843 - val\_accuracy: 0.9948 - val\_loss: 0.0221  
Epoch 132/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9054 - loss: 0.3496 - val\_accuracy: 0.9768 - val\_loss: 0.1839  
Epoch 133/200

49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.9228 - loss: 0.2703 - val\_accuracy: 1.0000 - val\_loss: 0.0029  
Epoch 134/200

49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.9115 - loss: 0.2843 - val\_accuracy: 0.9459 - val\_loss: 0.3315  
Epoch 135/200

49/49 \_\_\_\_\_ 17s 357ms/step - accuracy: 0.8854 - loss: 0.3631 - val\_accuracy: 0.9820 - val\_loss: 0.0759  
Epoch 136/200

49/49 \_\_\_\_\_ 18s 358ms/step - accuracy: 0.9129 - loss: 0.2653 - val\_accuracy: 0.9794 - val\_loss: 0.1044  
Epoch 137/200

49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9235 - loss: 0.2680 - val\_accuracy: 0.9897 - val\_loss: 0.0333  
Epoch 138/200

49/49 \_\_\_\_\_ 18s 371ms/step - accuracy: 0.8897 - loss: 0.3498 - val\_accuracy: 0.9794 - val\_loss: 0.0547  
Epoch 139/200

49/49 \_\_\_\_\_ 18s 371ms/step - accuracy: 0.9241 - loss: 0.2330 - val\_accuracy: 0.9433 - val\_loss: 0.5016  
Epoch 140/200

49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.9249 - loss: 0.2402 - val\_accuracy: 0.9227 - val\_loss: 0.6250  
Epoch 141/200

49/49 \_\_\_\_\_ 18s 358ms/step - accuracy: 0.9117 - loss: 0.2718 - val\_accuracy: 0.9510 - val\_loss: 0.5088  
Epoch 142/200

49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.9075 - loss: 0.3041 - val\_accuracy: 1.0000 - val\_loss: 0.0147  
Epoch 143/200

49/49 \_\_\_\_\_ 18s 362ms/step - accuracy: 0.8967 - loss: 0.3152 - val\_accuracy: 0.9871 - val\_loss: 0.0296  
Epoch 144/200

49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9078 - loss: 0.2538 - val\_accuracy: 0.9716 - val\_loss: 0.1290  
Epoch 145/200

49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.9035 - loss: 0.2724 - val\_accuracy: 0.9768 - val\_loss: 0.0658  
Epoch 146/200

49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9060 - loss: 0.2732 - val\_accuracy: 0.9716 - val\_loss: 0.1267  
Epoch 147/200

49/49 \_\_\_\_\_ 18s 360ms/step - accuracy: 0.9254 - loss: 0.2392 - val\_accuracy: 0.9871 - val\_loss: 0.0220  
Epoch 148/200

49/49 \_\_\_\_\_ 18s 366ms/step - accuracy: 0.8977 - loss: 0.3027 - val\_accuracy: 0.9562 - val\_loss: 0.2223  
Epoch 149/200

49/49 \_\_\_\_\_ 18s 360ms/step - accuracy: 0.9289 - loss:

0.2725 - val\_accuracy: 0.9485 - val\_loss: 0.3669  
Epoch 150/200  
49/49 \_\_\_\_\_ 18s 358ms/step - accuracy: 0.9223 - loss: 0.2288 - val\_accuracy: 0.9923 - val\_loss: 0.0217  
Epoch 151/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.8924 - loss: 0.3316 - val\_accuracy: 1.0000 - val\_loss: 0.0076  
Epoch 152/200  
49/49 \_\_\_\_\_ 18s 358ms/step - accuracy: 0.9256 - loss: 0.2392 - val\_accuracy: 0.9768 - val\_loss: 0.1598  
Epoch 153/200  
49/49 \_\_\_\_\_ 17s 354ms/step - accuracy: 0.9315 - loss: 0.2137 - val\_accuracy: 0.9794 - val\_loss: 0.0875  
Epoch 154/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.9147 - loss: 0.2321 - val\_accuracy: 0.9794 - val\_loss: 0.0548  
Epoch 155/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.8884 - loss: 0.3556 - val\_accuracy: 0.9330 - val\_loss: 0.5892  
Epoch 156/200  
49/49 \_\_\_\_\_ 18s 359ms/step - accuracy: 0.9205 - loss: 0.2652 - val\_accuracy: 0.9433 - val\_loss: 0.2727  
Epoch 157/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9114 - loss: 0.2597 - val\_accuracy: 0.9716 - val\_loss: 0.0795  
Epoch 158/200  
49/49 \_\_\_\_\_ 18s 357ms/step - accuracy: 0.9017 - loss: 0.3223 - val\_accuracy: 0.9820 - val\_loss: 0.2644  
Epoch 159/200  
49/49 \_\_\_\_\_ 18s 369ms/step - accuracy: 0.9052 - loss: 0.2728 - val\_accuracy: 0.9639 - val\_loss: 0.2694  
Epoch 160/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.9078 - loss: 0.3011 - val\_accuracy: 0.9742 - val\_loss: 0.2167  
Epoch 161/200  
49/49 \_\_\_\_\_ 17s 353ms/step - accuracy: 0.9119 - loss: 0.2435 - val\_accuracy: 0.9562 - val\_loss: 0.3107  
Epoch 162/200  
49/49 \_\_\_\_\_ 17s 356ms/step - accuracy: 0.9137 - loss: 0.2730 - val\_accuracy: 0.9691 - val\_loss: 0.2991  
Epoch 163/200  
49/49 \_\_\_\_\_ 17s 355ms/step - accuracy: 0.8981 - loss: 0.2893 - val\_accuracy: 0.9820 - val\_loss: 0.1145  
Epoch 164/200  
49/49 \_\_\_\_\_ 27s 546ms/step - accuracy: 0.9204 - loss: 0.2585 - val\_accuracy: 0.9356 - val\_loss: 0.5077  
Epoch 165/200  
49/49 \_\_\_\_\_ 42s 861ms/step - accuracy: 0.8931 - loss: 0.3987 - val\_accuracy: 0.9536 - val\_loss: 0.1747  
Epoch 166/200



49/49 \_\_\_\_\_ 36s 743ms/step - accuracy: 0.9204 - loss: 0.2347 - val\_accuracy: 0.9536 - val\_loss: 0.3349  
Epoch 167/200

49/49 \_\_\_\_\_ 38s 782ms/step - accuracy: 0.9423 - loss: 0.1808 - val\_accuracy: 0.9639 - val\_loss: 0.4547  
Epoch 168/200

49/49 \_\_\_\_\_ 37s 754ms/step - accuracy: 0.9367 - loss: 0.1804 - val\_accuracy: 0.9381 - val\_loss: 0.5335  
Epoch 169/200

49/49 \_\_\_\_\_ 35s 713ms/step - accuracy: 0.9188 - loss: 0.2452 - val\_accuracy: 0.9897 - val\_loss: 0.0372  
Epoch 170/200

49/49 \_\_\_\_\_ 44s 899ms/step - accuracy: 0.9181 - loss: 0.2440 - val\_accuracy: 0.9588 - val\_loss: 0.3276  
Epoch 171/200

49/49 \_\_\_\_\_ 42s 854ms/step - accuracy: 0.9147 - loss: 0.2964 - val\_accuracy: 0.9691 - val\_loss: 0.1067  
Epoch 172/200

49/49 \_\_\_\_\_ 41s 826ms/step - accuracy: 0.9185 - loss: 0.2887 - val\_accuracy: 0.9897 - val\_loss: 0.0352  
Epoch 173/200

49/49 \_\_\_\_\_ 38s 778ms/step - accuracy: 0.9247 - loss: 0.2582 - val\_accuracy: 0.9923 - val\_loss: 0.0286  
Epoch 174/200

49/49 \_\_\_\_\_ 37s 754ms/step - accuracy: 0.9261 - loss: 0.2392 - val\_accuracy: 0.9974 - val\_loss: 0.0099  
Epoch 175/200

49/49 \_\_\_\_\_ 36s 734ms/step - accuracy: 0.9253 - loss: 0.2377 - val\_accuracy: 0.8789 - val\_loss: 1.2572  
Epoch 176/200

49/49 \_\_\_\_\_ 38s 780ms/step - accuracy: 0.9401 - loss: 0.2017 - val\_accuracy: 0.9820 - val\_loss: 0.0888  
Epoch 177/200

49/49 \_\_\_\_\_ 37s 752ms/step - accuracy: 0.9395 - loss: 0.1776 - val\_accuracy: 0.9871 - val\_loss: 0.0704  
Epoch 178/200

49/49 \_\_\_\_\_ 35s 717ms/step - accuracy: 0.9187 - loss: 0.2321 - val\_accuracy: 0.9820 - val\_loss: 0.0934  
Epoch 179/200

49/49 \_\_\_\_\_ 35s 706ms/step - accuracy: 0.9189 - loss: 0.2597 - val\_accuracy: 0.9974 - val\_loss: 0.0282  
Epoch 180/200

49/49 \_\_\_\_\_ 34s 689ms/step - accuracy: 0.9341 - loss: 0.2191 - val\_accuracy: 0.9871 - val\_loss: 0.0773  
Epoch 181/200

49/49 \_\_\_\_\_ 35s 705ms/step - accuracy: 0.9203 - loss: 0.2256 - val\_accuracy: 0.9691 - val\_loss: 0.2415  
Epoch 182/200

49/49 \_\_\_\_\_ 33s 664ms/step - accuracy: 0.9340 - loss:



0.1910 - val\_accuracy: 0.9768 - val\_loss: 0.1829  
Epoch 183/200  
49/49 \_\_\_\_\_ 34s 685ms/step - accuracy: 0.9319 - loss: 0.2004 - val\_accuracy: 0.9330 - val\_loss: 0.7322  
Epoch 184/200  
49/49 \_\_\_\_\_ 35s 711ms/step - accuracy: 0.9162 - loss: 0.2704 - val\_accuracy: 0.9742 - val\_loss: 0.0856  
Epoch 185/200  
49/49 \_\_\_\_\_ 36s 742ms/step - accuracy: 0.9254 - loss: 0.2111 - val\_accuracy: 1.0000 - val\_loss: 0.0102  
Epoch 186/200  
49/49 \_\_\_\_\_ 36s 731ms/step - accuracy: 0.9273 - loss: 0.2148 - val\_accuracy: 0.9562 - val\_loss: 0.2245  
Epoch 187/200  
49/49 \_\_\_\_\_ 36s 726ms/step - accuracy: 0.9467 - loss: 0.1592 - val\_accuracy: 0.9691 - val\_loss: 0.2889  
Epoch 188/200  
49/49 \_\_\_\_\_ 33s 672ms/step - accuracy: 0.9484 - loss: 0.1601 - val\_accuracy: 0.9665 - val\_loss: 0.2294  
Epoch 189/200  
49/49 \_\_\_\_\_ 34s 698ms/step - accuracy: 0.9275 - loss: 0.1938 - val\_accuracy: 0.9948 - val\_loss: 0.0086  
Epoch 190/200  
49/49 \_\_\_\_\_ 33s 674ms/step - accuracy: 0.9168 - loss: 0.2633 - val\_accuracy: 0.9510 - val\_loss: 0.6461  
Epoch 191/200  
49/49 \_\_\_\_\_ 33s 670ms/step - accuracy: 0.8864 - loss: 0.3120 - val\_accuracy: 0.9665 - val\_loss: 0.3026  
Epoch 192/200  
49/49 \_\_\_\_\_ 34s 690ms/step - accuracy: 0.9327 - loss: 0.2127 - val\_accuracy: 0.9845 - val\_loss: 0.1167  
Epoch 193/200  
49/49 \_\_\_\_\_ 35s 704ms/step - accuracy: 0.9209 - loss: 0.2678 - val\_accuracy: 0.9897 - val\_loss: 0.0604  
Epoch 194/200  
49/49 \_\_\_\_\_ 34s 689ms/step - accuracy: 0.9136 - loss: 0.2622 - val\_accuracy: 0.9974 - val\_loss: 0.0110  
Epoch 195/200  
49/49 \_\_\_\_\_ 33s 674ms/step - accuracy: 0.9336 - loss: 0.1965 - val\_accuracy: 0.9794 - val\_loss: 0.2545  
Epoch 196/200  
49/49 \_\_\_\_\_ 33s 665ms/step - accuracy: 0.9374 - loss: 0.1973 - val\_accuracy: 0.9794 - val\_loss: 0.2142  
Epoch 197/200  
49/49 \_\_\_\_\_ 34s 696ms/step - accuracy: 0.9432 - loss: 0.1837 - val\_accuracy: 0.9845 - val\_loss: 0.0786  
Epoch 198/200  
49/49 \_\_\_\_\_ 25s 512ms/step - accuracy: 0.9332 - loss: 0.2226 - val\_accuracy: 0.9716 - val\_loss: 0.2337

Epoch 199/200

49/49 ————— 16s 335ms/step - accuracy: 0.9362 - loss: 0.2030 - val\_accuracy: 0.9794 - val\_loss: 0.2209

Epoch 200/200

49/49 ————— 17s 337ms/step - accuracy: 0.9600 - loss: 0.1611 - val\_accuracy: 0.9845 - val\_loss: 0.0438



Model saved as 'animal\_classifier.keras'

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import os

tf.config.set_visible_devices([], 'GPU')
print("Configured to use CPU only.")

MODEL_PATH = 'animal_classifier.keras'
IMAGE_PATH = "C:\\Users\\tanis\\Downloads\\bhalu.jpg"

IMAGE_SIZE = (128, 128)

CLASS_NAMES = sorted([
    'Bear', 'Bird', 'Cat', 'Cow', 'Deer', 'Dog', 'Dolphin',
    'Elephant',
    'Giraffe', 'Horse', 'Kangaroo', 'Lion', 'Panda', 'Tiger', 'Zebra'
])

model = tf.keras.models.load_model(MODEL_PATH)

def preprocess_image(image_path, target_size):
    img = tf.keras.preprocessing.image.load_img(image_path,
```

```

target_size=target_size)
    img_array = tf.keras.preprocessing.image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0)
    return img_array

test_image_array = preprocess_image(IMAGE_PATH, IMAGE_SIZE)

predictions = model.predict(test_image_array)

predicted_probabilities = predictions[0]
predicted_class_index = np.argmax(predicted_probabilities)
predicted_class_name = CLASS_NAMES[predicted_class_index]
confidence = predicted_probabilities[predicted_class_index] * 100

print(f"\nPredicted class: {predicted_class_name}")
print(f"Confidence: {confidence:.2f}%")

plt.figure(figsize=(6, 6))
display_img = tf.keras.preprocessing.image.load_img(IMAGE_PATH,
target_size=IMAGE_SIZE)
plt.imshow(display_img)
plt.title(f"Prediction: {predicted_class_name} ({confidence:.2f}%)")
plt.axis('off')
plt.show()

```

Configured to use CPU only.  
1/1 \_\_\_\_\_ 0s 162ms/step

Predicted class: Bear  
Confidence: 99.91%

Prediction: Bear (99.91%)

